

Response time: How long it takes to do a task

Throughput (吞吐量): Total work done per unit time

Relative Performance

■ Define Performance = $1/\text{Execution Time}$

■ "X is n time faster than Y"

$$\begin{aligned} \text{Performance}_X / \text{Performance}_Y \\ &= \text{Execution time}_Y / \text{Execution time}_X \\ &= n \end{aligned}$$

Execution Time

■ Elapsed time 运行时间


- ◆ Total response time, including all aspects
 - Processing, I/O, OS overhead, idle time
- ◆ Determines system performance

■ CPU time CPU时间

- ◆ Time spent processing a given job
 - Discounts I/O time, other jobs' shares
- ◆ Comprises user CPU time and system CPU time
- ◆ Different programs are affected differently by CPU and system performance

CPU Clocking

■ Operation of digital hardware governed by a constant-rate clock



■ Clock period: duration of a clock cycle

- ◆ e.g., $250\text{ps} = 0.25\text{ns} = 250 \times 10^{-12}\text{s}$

■ Clock frequency (rate): cycles per second

- ◆ e.g., $4.0\text{GHz} = 4000\text{MHz} = 4.0 \times 10^9\text{Hz}$

CPU Time

CPU Time = No. of Clock Cycles \times Clock Period

$$= \frac{\text{No. of Clock Cycles}}{\text{Clock Rate}}$$

■ Performance improved by

- ◆ Reducing number of clock cycles (cycle count)
- ◆ Increasing clock rate
- ◆ Hardware designer must often trade off clock rate against cycle count

Instruction Count and CPI

$$\text{Clock Cycles} = \text{Instruction Count} \times \text{Cycles per Instruction}$$

$$\text{CPU Time} = \text{Instruction Count} \times \text{CPI} \times \text{Clock Cycle Time}$$

$$= \frac{\text{Instruction Count} \times \text{CPI}}{\text{Clock Rate}}$$

- Instruction Count for a program
 - ◆ Determined by program, ISA and compiler
- Average cycles per instruction
 - ◆ Determined by CPU hardware
 - ◆ If different instructions have different CPI
 - Average CPI affected by instruction mix

- If different instruction classes take different numbers of cycles

$$\text{Clock Cycles} = \sum_{i=1}^n (\text{CPI}_i \times \text{Instruction Count}_i)$$

- Weighted average CPI

$$\text{CPI} = \frac{\text{Clock Cycles}}{\text{Instruction Count}} = \sum_{i=1}^n \left(\underbrace{\text{CPI}_i \times \frac{\text{Instruction Count}_i}{\text{Instruction Count}}}_{\text{Relative frequency}} \right)$$

Summary

$$\begin{aligned} \text{CPU Time} &= \frac{\text{Instructions}}{\text{Program}} \times \frac{\text{Clock cycles}}{\text{Instruction}} \times \frac{\text{Seconds}}{\text{Clock cycle}} \\ &= \text{IC} \times \text{CPI} \times T_c \end{aligned}$$


- Performance depends on
 - ◆ Algorithm: affects IC, possibly CPI
 - ◆ Programming language: affects IC, CPI
 - ◆ Compiler: affects IC, CPI
 - ◆ Instruction set architecture: affects IC, CPI, T_c

Energy Consumption

- Energy consumption = dynamic energy + static energy
 - ◆ Dynamic energy (energy spent when transistors switch from 0→1 1→0) is primary
 - ◆ Static energy is the energy cost when no transistor switches
- Energy for 0→1→0: $\text{Energy} \propto \text{Capacitive load} \times \text{Voltage}^2$
- Energy for 0→1 or 1→0: $\text{Energy} \propto 1/2 \times \text{Capacitive load} \times \text{Voltage}^2$
- Energy per second (power):

$$\text{Power} \propto 1/2 \times \text{Capacitive load} \times \text{Voltage}^2 \times \text{Frequency switched}$$

Power Trends



- In CMOS IC technology

$$\text{Power} \propto \frac{1}{2} \times \text{Capacitive load} \times \text{Voltage}^2 \times \text{Frequency}$$

× 23

5V → 1.5V

× 270

Multiprocessors

- Multicore microprocessors
 - ◆ More than one processor per chip
- Requires explicitly parallel programming
 - ◆ Compare with instruction level parallelism
 - Hardware executes multiple instructions at once
 - Hidden from the programmer
 - ◆ Hard to do
 - Programming for performance
 - Load balancing
 - Optimizing communication and synchronization

Benchmark Suites

- Each vendor announces a SPEC rating for their system
 - ◆ a measure of execution time for a fixed collection of programs
 - ◆ is a function of a specific CPU, memory system, IO system, operating system, compiler
 - ◆ enables easy comparison of different systems
- The key is coming up with a collection of relevant programs

SPEC CPU Benchmark

- Programs used to measure performance
 - ◆ Supposedly typical of actual workload
- Standard Performance Evaluation Cooperative (SPEC)
 - ◆ Develops benchmarks for CPU, I/O, Web, ...
- SPEC CPU2006
 - ◆ Elapsed time to execute a selection of programs
 - Negligible I/O, so focuses on CPU performance
 - ◆ Normalized relative to reference machine
 - ◆ Summarize as geometric mean of performance ratios
 - CINT2006 (integer) and CFP2006 (floating-point)

$$\sqrt[n]{\prod_{i=1}^n \text{Execution time ratio}_i}$$

SPEC Power Benchmark

- Power consumption of the server at different workload levels
 - ◆ Performance: ssj_ops (server side Java operations per second)
 - ◆ Power: Watts (Joules/sec)

$$\text{Overall ssj_ops per Watt} = \left(\sum_{i=0}^{10} \text{ssj_ops}_i \right) / \left(\sum_{i=0}^{10} \text{power}_i \right)$$

Amdahl's Law

- Architecture design is very **bottleneck-driven** – make the common case fast, do not waste resources on a component that has little impact on overall performance/power
- Amdahl's Law: performance improvements through an enhancement is limited by the **fraction of time** the enhancement comes into play

Processor : control , datapath